STAT 4100 Lecture Note

Week Four (Sep 26 & 28, 2022)

Zhiyang Zhou (zhiyang.zhou@umanitoba.ca, zhiyanggeezhou.github.io)

2022/Sep/28 21:37:10

Generating functions (con'd)

Moment generating function (con'd)

- Application
  - Characterizing distributions: $M_{\mathbf{X}}(\boldsymbol{t})$ and $M_{\mathbf{Y}}(\boldsymbol{t})$ are both well-defined and equal for all $\boldsymbol{t}$ in a neighborhood of $\mathbf{0} \Leftrightarrow \mathbf{X} \stackrel{d}{=} \mathbf{Y}$
    - Proofs for laws of large numbers and central limit theorems.
  - Computing moments
    - nth raw moment $\mu_n' = \mathrm{E}X^n = \sum_{k=0}^n \binom{n}{k} \mu_k (\mu_1')^{n-k}$
    - nth central moment $\mu_n = \mathrm{E}(X - \mathrm{E}X)^n = \sum_{k=0}^n \binom{n}{k} (-1)^{n-k} \mu_k' (\mu_1')^{n-k}$

Characteristic function

- For univariate X: $\phi_X(t) = \mathrm{E}\exp(itX)$ for all $t \in \mathbb{R}$
  - Fourier transform of f_X
  - Inverse: $f_X(x) = (2\pi)^{-1} \int_{\mathbb{R}} \phi_X(t) \exp(-itx) \mathrm{d}t$
  - $\mu_n' = \mathrm{E}X^n = (-i)^n \phi_X^{(n)}(0)$
- For Multivariate $\mathbf{X} = (X_1, \ldots, X_p)^\top$: $\phi_{\mathbf{X}}(\boldsymbol{t}) = \mathrm{E}\exp(i\boldsymbol{t}^\top \mathbf{X})$ for all $\boldsymbol{t} \in \mathbb{R}^p$
  - Fourier transform of $f_{\mathbf{X}}$
  - Inverse: $f_{\mathbf{X}}(\boldsymbol{x}) = (2\pi)^{-p} \int_{\mathbb{R}^p} \phi_{\mathbf{X}}(\boldsymbol{t}) \exp(-i\boldsymbol{t}^\top \boldsymbol{x}) \mathrm{d}\boldsymbol{t}$
- $\phi_{\mathbf{X}}(\boldsymbol{t}) = \phi_{\mathbf{Y}}(\boldsymbol{t})$ for all $\boldsymbol{t} \in \mathbb{R}^p \Leftrightarrow \mathbf{X} \stackrel{d}{=} \mathbf{Y}$

Example Lec6.2

- Find the characteristic functions of following distributions.
  - $\mathcal{N}(\mu, \sigma^2)$.
  - $\mathrm{MVN}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$.
  - Cauchy distribution: $f_X(x) = \{\pi(1 + x^2)\}^{-1}$, $x \in \mathbb{R}$.

Other generating functions

- Cumulant generating function
  - $K_X(t) = \ln M_X(t) = \sum_{n=0}^\infty \kappa_n t^n / n!$
  - $\kappa_n = K_X^{(n)}(0)$
- Probability-generating function
  - For discrete r.v. X taking values from $\{0, 1, \ldots\}$, $G(z) = \mathrm{E}t^X = \sum_{x=0}^\infty t^x p_X(x)$.
  - $p_X(n) = \Pr(X = n) = G^{(n)}(1)/n!$

Estimating equations

Parametric models

- A parametric model is a set of distributions indexed by unknown $\boldsymbol{\theta} \in \boldsymbol{\Theta} \subset \mathbb{R}^p$ with small or moderate p
  - Say $\{f(\cdot \mid \boldsymbol{\theta}) : \boldsymbol{\theta} \in \boldsymbol{\Theta} \subset \mathbb{R}^p\}$, where f is either a pdf or a pmf and $\boldsymbol{\Theta}$ is the set of all the possbile values of $\boldsymbol{\theta}$
- Believed that the true parameter (vector) $\boldsymbol{\theta}_0$ ($\in \boldsymbol{\Theta} \subset \mathbb{R}^p$) is fixed
  - Rather than making $\boldsymbol{\theta}_0$ random in the Bayesian philosophy

Exponential family (CB Sec 3.4)

- Original parameterization

$$f(x \mid \boldsymbol{\theta}) = h(x)c(\boldsymbol{\theta}) \exp\left\{\sum_{i=1}^{k} w_i(\boldsymbol{\theta}) t_i(x)\right\}$$

- Normal (CB Example 3.4.4):
  - $h(x) = \mathbf{1}_{\mathbb{R}}(x)$
  - $c(\mu, \sigma) = (2\pi\sigma^2)^{-1/2} \exp\{-\mu^2/(2\sigma^2)\}\mathbf{1}_{\mathbb{R}}(\mu)\mathbf{1}_{\mathbb{R}^+}(\sigma)$
  - $w_1(\mu, \sigma) = \sigma^{-2}\mathbf{1}_{\mathbb{R}^+}(\sigma)$ & $w_2(\mu, \sigma) = \mu\sigma^{-2}\mathbf{1}_{\mathbb{R}^+}(\sigma)$
  - $t_1(x) = -x^2/2$ & $t_2(x) = x$
- Binomial (CB Example 3.4.1):
  - $h(x) = \binom{n}{x}\mathbf{1}_{\{0,\ldots,n\}}(x)$
  - $c(p) = (1-p)^n \mathbf{1}_{(0,1)}(p)$
  - $w_1(p) = \ln\{p/(1-p)\}\mathbf{1}_{(0,1)}(p)$
  - $t_1(x) = x$
- Other special cases: gamma, beta, Poisson, negative binomial

Method of moments (MOM, CB Sec 7.2.1)

- Procedure
  1. Equate raw moments to their empirical counterparts.
  2. Solve the resulting simultaneous equations for $\boldsymbol{\theta} = (\theta_1, \ldots, \theta_p)$.
- Features
  - Easy implementation
  - Start point for more complex methods
  - No constraint
  - Not uniquely defined
  - No guarantee on optimality

Exercise Lec7.1

- Let $X_1, \ldots, X_n$ iid follow the following distributions. Find MOM estimators for (θ_1, θ_2).
  a. $N(\theta_1, \theta_2)$, $(\theta_1, \theta_2) \in \mathbb{R} \times \mathbb{R}^+$.
  b. $\text{Binom}(\theta_1, \theta_2)$ with pmf

$$p_X(x \mid \theta_1, \theta_2) = \binom{\theta_1}{x}\theta_2^x(1-\theta_2)^{\theta_1 - x}\mathbf{1}_{\{0,\ldots,\theta_1\}}(x), \quad (\theta_1, \theta_2) \in \mathbb{Z}^+ \times (0,1).$$

Exercise Lec7.2

- Let $X_1, \ldots, X_n$ iid follow pdf $f(x \mid \theta) = \theta x^{\theta-1}\mathbf{1}_{[0,1]}(x)$, $\theta > 0$.

a. Find an MOM estimator of θ.
b. Can we employ the second (raw) moment instead of the first one?

Maximum Likelihood Estimator (MLE, CB Sec 7.2.2)

- Likelihood function: $L : \boldsymbol{\Theta} \to \mathbb{R}$ such that, given $\boldsymbol{x}$ (a realization of $\mathbf{X}$),

$$L(\boldsymbol{\theta}) = L(\boldsymbol{\theta}; \boldsymbol{x}) = f_{\mathbf{X}}(\boldsymbol{x} \mid \boldsymbol{\theta}),$$

  where $f_{\mathbf{X}}$ is the joint pdf or pmf.

- For each $\boldsymbol{x}$, let $\hat{\boldsymbol{\theta}}(\boldsymbol{x})$ be the maximizer of $L(\boldsymbol{\theta}; \boldsymbol{x})$ (or log-likelihood $\ell(\boldsymbol{\theta}; \boldsymbol{x}) = \ln L(\boldsymbol{\theta}; \boldsymbol{x})$) with respect to $\boldsymbol{\theta}$ constrained in $\boldsymbol{\Theta}$, i.e.,

$$\hat{\boldsymbol{\theta}}(\boldsymbol{x}) = \arg\max_{\boldsymbol{\theta} \in \boldsymbol{\Theta}} L(\boldsymbol{\theta}; \boldsymbol{x}) = \arg\max_{\boldsymbol{\theta} \in \boldsymbol{\Theta}} \ell(\boldsymbol{\theta}; \boldsymbol{x}).$$

  Then the statistic $\hat{\boldsymbol{\theta}} = \hat{\boldsymbol{\theta}}(\mathbf{X})$ is the MLE for $\boldsymbol{\theta} \in \boldsymbol{\Theta}$.

- Invariance property of MLE (CB Thm 7.2.10): As long as $\hat{\boldsymbol{\theta}}$ is the MLE of $\boldsymbol{\theta}$, for ANY function g, the $g(\hat{\boldsymbol{\theta}})$ is teh MLE of $g(\boldsymbol{\theta})$.

- If ℓ is differetiable, the score funtion $\mathbf{S}$ is defined as its gradient

$$\mathbf{S}(\boldsymbol{\theta}) = \mathbf{S}(\boldsymbol{\theta}; \boldsymbol{x}) = \left[\frac{\partial}{\partial \theta_1} \ell(\boldsymbol{\theta}; \boldsymbol{x}), \ldots, \frac{\partial}{\partial \theta_p} \ell(\boldsymbol{\theta}; \boldsymbol{x})\right]^\top.$$

- If ℓ is twice differentiable, we have hessian of $\ell(\boldsymbol{\theta}; \boldsymbol{x})$

$$\mathbf{H}(\boldsymbol{\theta}) = \mathbf{H}(\boldsymbol{\theta}; \boldsymbol{x}) = \left[\frac{\partial^2}{\partial \theta_i \partial \theta_j} \ell(\boldsymbol{\theta}; \boldsymbol{x})\right]_{p \times p}.$$

- Maximizing twice-differentiable ℓ
  1. Find out stationary points, i.e., solutions to simultaneous equations $S(\boldsymbol{\theta}) = \mathbf{0}$
  2. Screen out (interior) local maximizers, i.e., stationary points with negative definite Hessian matrix
  3. Determine the global maximizer within $\boldsymbol{\Theta}$: by comparing values of likelihood (or log-likelihood) evaluated at local maximizers and boundary points of $\boldsymbol{\Theta}$

Exercise Lec7.3

- Suppose $X_1, \ldots, X_n$ are iid as the following distributions. Find MLEs for corresponding parameters.
  a. $N(\mu, \sigma^2)$, $(\mu, \sigma) \in \mathbb{R} \times \mathbb{R}^+$.
  b. Bernoulli(θ): $p(x \mid \theta) = \theta^x (1-\theta)^{1-x} \mathbf{1}_{\{0,1\}}(x)$, $\theta \in [0, 1/2]$.
  c. Two-parameter exponential: $f(x \mid \alpha, \beta) = \beta^{-1} \exp\{-(x-\alpha)/\beta\} \mathbf{1}_{(\alpha,\infty)}(x)$, $(\alpha, \beta) \in \mathbb{R} \times \mathbb{R}^+$.

Other examples of estimating equations

- Least-squares estimator
- Generalized estimating equations (GEE)
- M-estimator