

# STAT 3690 Lecture Note

## Part X: Clustering

Zhiyang Zhou (zhiyang.zhou@umanitoba.ca, zhiyanggeezhou.github.io)

2023/Apr/09 20:41:23

---

### Clustering

- Problem: given observations  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^p$  group the observations into  $K$  populations
    - Unknown  $K$
    - Unsupervised: no label/training data
  - Why
    - Summarize a representation of the full data set
    - Exploration for structure of the data
    - Checking the validity of pre-existing group assignments
    - Assistance for prediction: sometimes clustering prior to prediction
  - Find a function  $C$  such that
    - $C(i) = k$ : assign the  $i$ th subject to group  $k$
- 

### $K$ -means

- Within-cluster scatter

$$W(K) = \frac{1}{2} \sum_{k=1}^K \frac{1}{n_k} \sum_{i: C(i)=k} \sum_{j: C(j)=k} \|\mathbf{x}_i - \mathbf{x}_j\|_2^2 = \sum_{k=1}^K \sum_{i: C(i)=k} \|\mathbf{x}_i - \bar{\mathbf{x}}_k\|_2^2$$

- $\|\mathbf{x}_i - \mathbf{x}_j\|_2$ : the Euclidean distance between  $\mathbf{x}_i$  and  $\mathbf{x}_j$
  - $\bar{\mathbf{x}}_k = n_k^{-1} \sum_{i: C(i)=k} \mathbf{x}_i$
  - Smaller  $W(K)$  is better
- (Approximately) minimizing the within-cluster scatter

$$\min_C W(K) = \min_{C, \mathbf{c}_1, \dots, \mathbf{c}_K} \sum_{k=1}^K \sum_{i: C(i)=k} \|\mathbf{x}_i - \mathbf{c}_k\|_2^2$$

- Implementation:
    1. Specify  $K$  and start with an initial guess for  $\mathbf{c}_1, \dots, \mathbf{c}_K$ , then repeat
      - a. Labeling each point based the closest center: for each  $i$ , put  $\mathbf{x}_i$  to the  $k$ th cluster such that  $\mathbf{c}_k$  is closest to  $\mathbf{x}_i$
      - b. Replacing each center by the average of points in its cluster: for each  $k$ , take  $\mathbf{c}_k = \bar{\mathbf{x}}_k$
    2. Terminate when  $W(K)$  doesn't change
-

- Comments
  - Always converge via the expectation maximization (EM) algorithm
  - No guarantee to lead to the smallest  $W$
  - Depend on  $K$  and initial cluster centers
    - \* Typically run  $K$ -means multiple times and pick up the result with the smallest  $W$
- Determining  $K$ 
  - Between-cluster variation

$$B(K) = \sum_{k=1}^K n_k \|\bar{\mathbf{x}}_k - \bar{\mathbf{x}}\|_2^2$$

- \*  $\bar{\mathbf{x}} = n^{-1} \sum_{i=1}^n \mathbf{x}_i$
- CH index (Caliński & Harabasz (1974), *Communications in Statistics*, 3:1–27)

$$\text{CH}(K) = \frac{B(K)/(K-1)}{W(K)/(n-K)}$$

- To choose  $K$  as the maximizer of  $\text{CH}(K)$ , i.e.,

$$\hat{K} = \arg \max_{K \in \{2, \dots, K_{\max}\}} \text{CH}(K)$$

- Example 10.1 ( $K$ -means for iris)

```
set.seed(1)
options(digits = 4)
x = iris[, !(names(iris) %in% c('Species'))]
y = (iris$Species == unique(iris$Species)[1]) +
  2*(iris$Species == unique(iris$Species)[2]) +
  3*(iris$Species == unique(iris$Species)[3])

## K-means via the first two PCs of x (to facilitate the visualization)
decomp = prcomp(x)
s = 2
PCscores = decomp$x[,1:s]
K = 3; cols = c("red", "darkgreen", "blue")
km = kmeans(PCscores, centers=K, nstart=1, algorithm="Lloyd", iter.max = 100)
# clustering plot with centers
plot(PCscores, col=cols[km$cluster])
points(km$centers, pch=19, cex=2, col=cols)
# comparison with true groups
par(mfrow=c(1,2))
plot(PCscores, col=cols[km$cluster], main="K-means")
plot(PCscores, col=cols[y], main="True")

## determine K for K-means via original x
set.seed(3690)
Ks = 2:20
Ws = numeric(length(Ks))
Bs = numeric(length(Ks))
CHs = numeric(length(Ks))
for(l in 1:length(Ks)){
  km = kmeans(x, centers=Ks[l], nstart=1, algorithm="Lloyd", iter.max = 100)
  Ws[l] = km$tot.withinss
  Bs[l] = sum(km$size * rowSums(sweep(km$centers, 2, colMeans(PCscores))^2))
  CHs[l] = (Bs[l]/(Ks[l]-1))/(Ws[l]/(nrow(PCscores)-Ks[l]))
}
```

```

}
plot(Ks, CHs,
     type="b", pch = 19,
     xlab="Number of clusters K",
     ylab="CH index")

```

- Color quantization/vector quantization (an application of  $K$ -means to image compression)
  - Basic idea: compress images by reducing the color palette of an image to  $K$  colors
  - Implementation:
    1. Let  $x_i$  be the quantified color of the  $i$ th pixel to be
    2. Initiate with  $K$  colors, say  $c_1, \dots, c_K$ , and then repeat the following steps until convergence
      - a. Classifying the  $i$ th pixel into the  $k$ th cluster if  $x_i$  is closest to  $c_k$  AND replacing  $x_i$  with  $c_k$
      - b. Updating  $c_k$  by  $\sum_{i \in \text{the } k\text{th cluster}} x_i / n_k$  with  $n_k$  as the size of the  $k$ th cluster

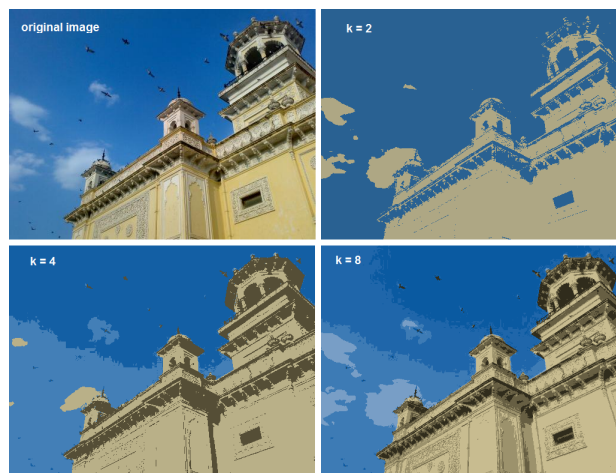


Figure 1: Image compression with  $K$ -means clustering (<http://opencvpython.blogspot.com/2012/12/k-means-clustering-2-working-with-scipy.html>)

- Example 10.2 (Color quantization for hand-written digits)

```

options(digits = 4)
mnist = dslabs::read_mnist()

# The 3690th image in the training set
i0 = 3690
image(z = matrix(mnist$train$images[i0,], ncol = 28),
      col = gray.colors(256, start = 0, rev = TRUE), axes = FALSE, main = "True")

# Shrink the number of colors to be 2
set.seed(1)
K = 2
km = kmeans(mnist$train$images[i0,], centers=K, nstart=1, algorithm="Lloyd", iter.max = 100)
image(z = matrix(km$cluster, ncol = 28),
      col = gray.colors(K, start = range(km$centers)[1]/255, end = range(km$centers)[2]/255, rev = T),
      axes = FALSE, main = "Compressed")

```

## Hierarchical clustering

- A simple example
  - Step 1:  $\{1\}, \{2\}, \{3\}, \{4\}, \{5\}, \{6\}, \{7\}$ ;
  - Step 2:  $\{1\}, \{2, 3\}, \{4\}, \{5\}, \{6\}, \{7\}$ ;
  - Step 3:  $\{1, 7\}, \{2, 3\}, \{4\}, \{5\}, \{6\}$ ;
  - Step 4:  $\{1, 7\}, \{2, 3\}, \{4, 5\}, \{6\}$ ;
  - Step 5:  $\{1, 7\}, \{2, 3, 6\}, \{4, 5\}$ ;
  - Step 6:  $\{1, 7\}, \{2, 3, 4, 5, 6\}$ ;
  - Step 7:  $\{1, 2, 3, 4, 5, 6, 7\}$ .
- Dendrogram: a tree displaying a hierarchical sequence of clustering assignments
  - Node representing a group
    - \* Leaf node representing a singleton (i.e., a group containing a single data point)
    - \* Root node representing the group containing all the data points
    - \* Internal node: has two children nodes, representing the the groups that were merged to form it
  - Height: draw each internal node at a height proportional to the dissimilarity between its two children nodes (if fix the leaf nodes at height zero)
- Distances
  - Dissimilarity  $d_{ij}$ : (Euclidean) distance between  $\mathbf{x}_i$  and  $\mathbf{x}_j$
  - Linkage: distance between groups  $G$  and  $H$ 
    - \* Options
      - Single linkage
$$d_{\text{single}}(G, H) = \min_{i \in G, j \in H} d_{ij}$$
      - Complete linkage
$$d_{\text{complete}}(G, H) = \max_{i \in G, j \in H} d_{ij}$$
      - Average linkage
$$d_{\text{average}}(G, H) = \frac{1}{n_G n_H} \sum_{i \in G, j \in H} d_{ij}$$
      - Centroid linkage
$$d_{\text{centroid}}(G, H) = \|\bar{\mathbf{x}}_G - \bar{\mathbf{x}}_H\|_2$$
      - Minimax linkage
$$d_{\text{minimax}}(G, H) = \min_{i \in G \cup H} \max_{j \in G \cup H} d_{ij}$$
    - \* Situation-dependent

- 
- Example 10.3 (hierarchical clustering for `iris`)
- 

## Modern alternatives

- Density-based spatial clustering of applications with noise (DBSCAN, M. Ester, H. Kriegel, J. Sander, X. Xu (1996), *Proceedings of the Second International Conference on Knowledge Discovery and Data Mining (KDD)*.)
- Uniform manifold approximation and projection for dimension reduction (UMAP, L. McInnes & J. Healy (2018), arXiv:1802.03426)