Week Four (Jan 30, Feb 1, & 3, 2023)

Zhiyang Zhou (zhiyang.zhou@umanitoba.ca, zhiyanggeezhou.github.io)

2023/Feb/01 10:48:53

Multivariate normal (MVN) distribution (con'd, J&W Sec 4.2)

Definition

- Standard MVN
  - $\boldsymbol{Z} = [Z_1, \ldots, Z_p]^\top \sim \mathrm{MVN}_p(\boldsymbol{0}, \mathbf{I}) \Leftrightarrow Z_1, \ldots, Z_p \overset{\mathrm{iid}}{\sim} \mathcal{N}(0, 1)$
  - pdf
  $$f_{\boldsymbol{Z}}(\boldsymbol{z}) = (2\pi)^{-p/2} \exp(-\boldsymbol{z}^\top \boldsymbol{z}/2) \cdot \mathbf{1}_{\mathbb{R}^p}(\boldsymbol{z})$$

- General MVN
  - $\boldsymbol{X} = [X_1, \ldots, X_p]^\top \sim \mathrm{MVN}_p(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \Leftrightarrow$ there exists $\boldsymbol{\mu} \in \mathbb{R}^p$, $\mathbf{A} \in \mathbb{R}^{p \times p}$ and $\boldsymbol{Z} \sim \mathrm{MVN}_p(\boldsymbol{0}, \mathbf{I})$ such that $\boldsymbol{X} = \mathbf{A}\boldsymbol{Z} + \boldsymbol{\mu}$ and $\boldsymbol{\Sigma} = \mathbf{A}\mathbf{A}^\top$
    * Limited to non-degenerate cases, i.e., invertible $\mathbf{A}$ ($\Leftrightarrow \boldsymbol{\Sigma} > 0$)
  - pdf
  $$f_{\boldsymbol{X}}(\boldsymbol{x}) = (2\pi)^{-p/2} (\det \boldsymbol{\Sigma})^{-1/2} \exp\{-(\boldsymbol{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\boldsymbol{x} - \boldsymbol{\mu})/2\} \cdot \mathbf{1}_{\mathbb{R}^p}(\boldsymbol{x})$$

- Exercise: Density of $\mathrm{MVN}_2(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ evaluated at $(4, 7)$, where

$$\boldsymbol{\mu} = [3, 6]^\top, \quad \boldsymbol{\Sigma} = \left[\begin{array}{cc} 10 & 2 \\ 2 & 5 \end{array}\right].$$

```r
options(digits = 4)
(Mu = matrix(c(3, 6), ncol = 1, nrow = 2))
(Sigma = matrix(c(10, 2 ,2, 5), ncol = 2, nrow = 2))
(x = c(4,7))
# Method 1: following the pdf
(2*pi)^{-length(Mu)/2}*det(Sigma)^{-.5}*exp(-drop(t(x-Mu)%*%solve(Sigma)%*%(x-Mu))/2)
# Method 2: via mvtnorm::dmvnorm()
mvtnorm::dmvnorm(x, mean = Mu, sigma = Sigma)
```

Properties of MVN

- $\boldsymbol{X}$ is of MVN $\Leftrightarrow a^\top \boldsymbol{X}$ is normally distributed for ALL non-zero $a \in \mathbb{R}^p$.
  - Warning: marginal normals do not imply the joint normal.
- If $\boldsymbol{X} \sim \mathrm{MVN}_p(\boldsymbol{\mu}, \boldsymbol{\Sigma})$, then $\mathbf{A}\boldsymbol{X} + \boldsymbol{b} \sim \mathrm{MVN}_q(\mathbf{A}\boldsymbol{\mu} + \boldsymbol{b}, \mathbf{A}\boldsymbol{\Sigma}\mathbf{A}^\top)$ for $\mathbf{A} \in \mathbb{R}^{q \times p}$ of full-row-rank. Specifically, if $\boldsymbol{X} \sim \mathrm{MVN}_p(\boldsymbol{\mu}, \boldsymbol{\Sigma})$, then
  - $\boldsymbol{\Sigma}^{-1/2}(\boldsymbol{X} - \boldsymbol{\mu}) \sim \mathrm{MVN}_p(\boldsymbol{0}, \mathbf{I})$ AND

– (Stochastic representation of MVN) there is $\boldsymbol{Z} \sim \text{MVN}_p(\boldsymbol{0}, \mathbf{I})$ such that $\boldsymbol{X} = \boldsymbol{\Sigma}^{1/2}\boldsymbol{Z} + \boldsymbol{\mu}$.
- $(\boldsymbol{X} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\boldsymbol{X} - \boldsymbol{\mu}) \sim \chi^2(p)$ if $\boldsymbol{X} \sim \text{MVN}_p(\boldsymbol{\mu}, \boldsymbol{\Sigma})$.

- Exercise: Generate six iid samples following bivariate normal $\text{MVN}_2(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ with

$$\boldsymbol{\mu} = [3, 6]^\top, \quad \boldsymbol{\Sigma} = \begin{bmatrix} 10 & 2 \\ 2 & 5 \end{bmatrix}.$$

```r
options(digits = 4)
set.seed(1)
(Mu = matrix(c(3, 6), ncol = 1, nrow = 2))
(Sigma = matrix(c(10, 2 ,2, 5), ncol = 2, nrow = 2))
n = 10
# Method 1: following the stochastic representation
sample1 = matrix(0, nrow = n, ncol = length(Mu))
for (i in 1:n) {
  sample1[i, ] = t(
    expm::sqrtm(Sigma) %*%
    matrix(rnorm(length(Mu)), nrow = length(Mu), ncol = 1) +
    Mu
 )
}
sample1
# Method 2: via MASS::mvrnorm()
(sample2 = MASS::mvrnorm(n, Mu, Sigma))
```

- Exercise: Suppose $X_1 \sim \mathcal{N}(0, 1)$. In the following two cases, verify that $X_2 \sim \mathcal{N}(0, 1)$ as well. Does $\boldsymbol{X} = [X_1, X_2]^\top$ follow an MVN in both cases?
    a. $X_2 = -X_1$;
    b. $X_2 = (2Y - 1)X_1$, where $Y \sim \text{Ber}(p)$ and $Y \perp\!\!\!\perp X_1$.

```r
options(digits = 4)
set.seed(1)
xsize = 1e4L
x1 = rnorm(xsize)
# case a
x2 = -x1
plot3D::hist3D(z=table(cut(x1, 100), cut(x2, 100)), border = "black") # 3d histogram of (x1, x2)
plot3D::image2D(z=table(cut(x1, 100), cut(x2, 100)), border = "black") # plot the support of joint pdf
# case b
Y = rbinom(n = xsize, 1, .3)
x2 = (2 * Y - 1) * x1
plot3D::hist3D(z=table(cut(x1, 100), cut(x2, 100)), border = "black") # 3d histogram of (x1, x2)
plot3D::image2D(z=table(cut(x1, 100), cut(x2, 100)), border = "black") # plot the support of joint pdf
```

Marginal and conditional MVN

- If $\boldsymbol{X} \sim \mathrm{MVN}_p(\boldsymbol{\mu}, \boldsymbol{\Sigma})$, where

$$\boldsymbol{X} = \begin{bmatrix} \boldsymbol{X}_1 \\ \boldsymbol{X}_2 \end{bmatrix}, \quad \boldsymbol{\mu} = \begin{bmatrix} \boldsymbol{\mu}_1 \\ \boldsymbol{\mu}_2 \end{bmatrix} \quad \text{and} \quad \boldsymbol{\Sigma} = \begin{bmatrix} \boldsymbol{\Sigma}_{11} & \boldsymbol{\Sigma}_{12} \\ \boldsymbol{\Sigma}_{21} & \boldsymbol{\Sigma}_{22} \end{bmatrix}$$

  with
    - random p_i-vector $\boldsymbol{X}_i$, $i = 1, 2$,
    - p_i-vector $\boldsymbol{\mu}_i$, $i = 1, 2$,
    - $p_i \times p_i$ matrix $\boldsymbol{\Sigma}_{ii} > 0$, $i = 1, 2$,
- then
    - (Marginals of MVN are still MVN) $\boldsymbol{X}_i \sim \mathrm{MVN}_{p_i}(\boldsymbol{\mu}_i, \boldsymbol{\Sigma}_{ii})$
    - $\boldsymbol{X}_i \mid \boldsymbol{X}_j = \boldsymbol{x}_j \sim \mathrm{MVN}_{p_i}(\boldsymbol{\mu}_{i|j}, \boldsymbol{\Sigma}_{i|j})$
        * $\boldsymbol{\mu}_{i|j} = \boldsymbol{\mu}_i + \boldsymbol{\Sigma}_{ij}\boldsymbol{\Sigma}_{jj}^{-1}(\boldsymbol{x}_j - \boldsymbol{\mu}_j)$
        * $\boldsymbol{\Sigma}_{i|j} = \boldsymbol{\Sigma}_{ii} - \boldsymbol{\Sigma}_{ij}\boldsymbol{\Sigma}_{jj}^{-1}\boldsymbol{\Sigma}_{ji}$
    - $\boldsymbol{X}_1 \perp\!\!\!\perp \boldsymbol{X}_2 \Leftrightarrow \boldsymbol{\Sigma}_{12} = \boldsymbol{0}$
        * Warning: the prerequisite for this equivalence is the joint normal of $\boldsymbol{X}_1$ and $\boldsymbol{X}_2$.

- Exercise: The argument $\boldsymbol{X}_1 \perp\!\!\!\perp \boldsymbol{X}_2 \Leftrightarrow \boldsymbol{\Sigma}_{12} = 0$ is based on $[\boldsymbol{X}_1^\top, \boldsymbol{X}_2^\top]^\top \sim \mathrm{MVN}$. That is, if $\boldsymbol{X}_1$ and $\boldsymbol{X}_2$ are both MVN BUT they are not jointly normal, the zero $\boldsymbol{\Sigma}_{12}$ doesn't suffice for the independence between $\boldsymbol{X}_1$ and $\boldsymbol{X}_2$. Recall the case b. in the previous exercise: $X_1 \sim \mathcal{N}(0,1)$ and $X_2 = (2Y-1)X_1$, where $Y \sim \mathrm{Ber}(p)$ and $Y \perp\!\!\!\perp X_1$. Verify that X_1 and X_2 are not independent of each other. (Hint: assume the independence and then check the support of $[X_1, X_2]^\top$.)

Checking normality (J&W Sec 4.6)

- Checking the univariate marginal distributions
    - Normal Q-Q plot
        * qqnorm(); car::qqPlot()
    - Univariate normality test
        * shapiro.test(); nortest::ad.test()
- Testing the multivariate normality
    - MVN::mvn()
- Checking the quadratic form
    - χ^2 Q-Q plot
        * $D_i^2 = (\boldsymbol{X}_i - \bar{\boldsymbol{X}})^\top \mathbf{S}^{-1}(\boldsymbol{X}_i - \bar{\boldsymbol{X}}) \approx \chi^2(p)$ if $\boldsymbol{X}_i \overset{\text{iid}}{\sim} \mathrm{MVN}_p(\boldsymbol{\mu}, \boldsymbol{\Sigma})$
        * qqplot(); car::qqPlot()

Detecting outliers (J&W Sec 4.7)

- Scatter plot of standardized values

- Check the points farthest from the origin in χ^2 Q-Q plot

Improving normality (J&W Sec 4.8)

- (Original) Box-Cox (power) transformation: transform positive x into

$$x^* = \begin{cases} (x^\lambda - 1)/\lambda & \lambda \neq 0 \\ \ln(x) & \lambda = 0 \end{cases}$$

  with λ selected with certain criterion

- – If $x \leq 0$, change it to be positive first.
- – See J. Tukey (1977). *Exploratory Data Analysis.* Boston: Addison-Wesley.
- Multivariate Box-Cox transformation

Maximum likelihood (ML) estimation of $\boldsymbol{\mu}$ and $\boldsymbol{\Sigma}$ (J&W Sec 4.3)

- Sample: $\boldsymbol{X}_1, \ldots, \boldsymbol{X}_n \overset{\text{iid}}{\sim} \text{MVN}_p(\boldsymbol{\mu}, \boldsymbol{\Sigma})$, $n > p$

- Parameter space: $\boldsymbol{\Theta} = \{(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \mid \boldsymbol{\mu} \in \mathbb{R}^p, \boldsymbol{\Sigma} \in \mathbb{R}^{p \times p}, \boldsymbol{\Sigma} > 0\}$

- Likelihood function

$$L(\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \prod_{i=1}^n \left[\frac{1}{\sqrt{(2\pi)^p \det(\boldsymbol{\Sigma})}} \exp\left\{ -\frac{1}{2}(\boldsymbol{X}_i - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\boldsymbol{X}_i - \boldsymbol{\mu}) \right\} \right]$$
$$= \frac{1}{\sqrt{(2\pi)^{np}\{\det(\boldsymbol{\Sigma})\}^n}} \exp\left\{ -\frac{1}{2}\sum_{i=1}^n (\boldsymbol{X}_i - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\boldsymbol{X}_i - \boldsymbol{\mu}) \right\}$$

- Log likelihood

$$\ell(\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \ln\{L(\boldsymbol{\mu}, \boldsymbol{\Sigma})\} = -\frac{np}{2}\ln(2\pi) - \frac{n}{2}\ln\{\det(\boldsymbol{\Sigma})\} - \frac{1}{2}\sum_{i=1}^n (\boldsymbol{X}_i - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\boldsymbol{X}_i - \boldsymbol{\mu})$$

- ML estimator

$$(\hat{\boldsymbol{\mu}}_{\text{ML}}, \widehat{\boldsymbol{\Sigma}}_{\text{ML}}) = \arg\max_{(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \in \boldsymbol{\Theta}} \ell(\boldsymbol{\mu}, \boldsymbol{\Sigma}) = (\bar{\boldsymbol{X}}, \frac{n-1}{n}\mathbf{S})$$

- – Consistency: $(\hat{\boldsymbol{\mu}}_{\text{ML}}, \widehat{\boldsymbol{\Sigma}}_{\text{ML}})$ approaches $(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ (in certain sense) as $n \to \infty$
- – Efficiency: the covariance matrix of $(\hat{\boldsymbol{\mu}}_{\text{ML}}, \widehat{\boldsymbol{\Sigma}}_{\text{ML}})$ is approximately optimal (in certain sense) as $n \to \infty$
- – Invariance: For any function g, the ML estimator of $g(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ is $g(\hat{\boldsymbol{\mu}}_{\text{ML}}, \widehat{\boldsymbol{\Sigma}}_{\text{ML}})$.

Sampling distributions of $\bar{X}$ and S (J&W Sec 4.4)

- Recall the univariate case
  - – $X_1, \ldots, X_n \overset{\text{iid}}{\sim} \mathcal{N}(\mu, \sigma^2)$
  - – $S^2 \perp\!\!\!\perp \bar{X}$
    - ∗ Sample variance $S^2 = (n-1)^{-1}\sum_{i=1}^n (X_i - \bar{X})^2$
  - – $\sqrt{n}(\bar{X} - \mu)/\sigma \sim \mathcal{N}(0, 1)$
  - – $(n-1)S^2/\sigma^2 \sim \chi^2(n-1)$
  - – $\sqrt{n}(\bar{X} - \mu)/S \sim t(n-1)$
- The multivariate case
  - – $\boldsymbol{X}_1, \ldots, \boldsymbol{X}_n \overset{\text{iid}}{\sim} \text{MVN}_p(\boldsymbol{\mu}, \boldsymbol{\Sigma})$, $n > p$
  - – $\mathbf{S} \perp\!\!\!\perp \bar{\boldsymbol{X}}$, i.e., $\widehat{\boldsymbol{\Sigma}}_{\text{ML}} \perp\!\!\!\perp \hat{\boldsymbol{\mu}}_{\text{ML}}$
  - – $\sqrt{n}\boldsymbol{\Sigma}^{-1/2}(\bar{\boldsymbol{X}} - \boldsymbol{\mu}) \sim \text{MVN}_p(\mathbf{0}, \mathbf{I})$
  - – $(n-1)\mathbf{S} = n\widehat{\boldsymbol{\Sigma}}_{\text{ML}} \sim W_p(\boldsymbol{\Sigma}, n-1)$
  - – $n(\bar{\boldsymbol{X}} - \boldsymbol{\mu})^\top \mathbf{S}^{-1}(\bar{\boldsymbol{X}} - \boldsymbol{\mu}) \sim$ Hotelling's $T^2(p, n-1)$

- Wishart distribution
  - – Def: $W_p(\boldsymbol{\Sigma}, n)$ is the distribution of $\sum_{i=1}^n \boldsymbol{Y}_i \boldsymbol{Y}_i^\top$ with $\boldsymbol{Y}_1, \ldots, \boldsymbol{Y}_n \overset{\text{iid}}{\sim} \text{MVN}_p(\mathbf{0}, \boldsymbol{\Sigma})$
    - ∗ A generalization of χ^2-distribution: $W_p(\boldsymbol{\Sigma}, n) = \chi^2(n)$ if $p = \boldsymbol{\Sigma} = 1$

- Propoties
  * $\mathbf{AA}^\top > 0$ and $\mathbf{W} \sim W_p(\mathbf{\Sigma}, n) \Rightarrow \mathbf{AWA}^\top \sim W_p(\mathbf{A\Sigma A}^\top, n)$
  * $\mathbf{W}_i \stackrel{\text{iid}}{\sim} W_p(\mathbf{\Sigma}, n_i) \Rightarrow \mathbf{W}_1 + \mathbf{W}_2 \sim W_p(\mathbf{\Sigma}, n_1 + n_2)$
  * $\mathbf{W}_1 \perp\!\!\!\perp \mathbf{W}_2$, $\mathbf{W}_1 + \mathbf{W}_2 \sim W_p(\mathbf{\Sigma}, n)$ and $\mathbf{W}_1 \sim W_p(\mathbf{\Sigma}, n_1) \Rightarrow \mathbf{W}_2 \sim W_p(\mathbf{\Sigma}, n - n_1)$
  * $\mathbf{W} \sim W_p(\mathbf{\Sigma}, n)$ and $\boldsymbol{a} \in \mathbb{R}^p \Rightarrow$

  $$\frac{\boldsymbol{a}^\top \mathbf{W} \boldsymbol{a}}{\boldsymbol{a}^\top \mathbf{\Sigma} \boldsymbol{a}} \sim \chi^2(n)$$

  * $\mathbf{W} \sim W_p(\mathbf{\Sigma}, n)$, $\boldsymbol{a} \in \mathbb{R}^p$ and $n \geq p \Rightarrow$

  $$\frac{\boldsymbol{a}^\top \mathbf{\Sigma}^{-1} \boldsymbol{a}}{\boldsymbol{a}^\top \mathbf{W}^{-1} \boldsymbol{a}} \sim \chi^2(n - p + 1)$$

  * $\mathbf{W} \sim W_p(\mathbf{\Sigma}, n) \Rightarrow$

  $$\operatorname{tr}(\mathbf{\Sigma}^{-1} \mathbf{W}) \sim \chi^2(np)$$

- Hotelling's T^2 distribution
  - A generalization of (Student's) t-distribution
  - If $\boldsymbol{X} \sim \text{MVN}_p(\mathbf{0}, \mathbf{I})$ and $\mathbf{W} \sim W_p(\mathbf{I}, n)$, then

  $$\boldsymbol{X}^\top \mathbf{W}^{-1} \boldsymbol{X} \sim T^2(p, n)$$

  - $Y \sim T^2(p, n) \Leftrightarrow \frac{n - p + 1}{np} Y \sim F(p, n - p + 1)$

- Wilk's lambda distribution
  - Wilks's lambda is to Hotelling's T^2 as F distribution is to Student's t in univariate statistics.
  - Given independent $\mathbf{W}_1 \sim W_p(\mathbf{\Sigma}, n_1)$ and $\mathbf{W}_2 \sim W_p(\mathbf{\Sigma}, n_2)$ with $n_1 \geq p$,

  $$\Lambda = \frac{\det(\mathbf{W}_1)}{\det(\mathbf{W}_1 + \mathbf{W}_2)} = \frac{1}{\det(\mathbf{I} + \mathbf{W}_1^{-1} \mathbf{W}_2)} \sim \Lambda(p, n_1, n_2)$$

    * Resort to an approximation in computation: $\{(p - n_2 + 1)/2 - n_1\} \ln \Lambda(p, n_1, n_2) \approx \chi^2(n_2 p)$

Inference on $\boldsymbol{\mu}$

Hypothesis testing

- Is it a squirrel?

Figure 1: Squirrel (Photograph by the Lacoste Garden Centre)

Figure 2: Flying Squirrel (Photograph by Joel Sartore)

Figure 3: Flying Squirrel (Photograph by Alex Badyaev)

- Null and alternative hypotheses, say H_0 and H_1, resp.

- Name of approach

- Test statistic (not unique) and corresponding level α rejection region R_α

  - $\Pr(\text{test statistic} \in R_\alpha \mid H_0) \leq \alpha$
  - Reject H_0 if the value of test statistic $\in R_\alpha$
    * Type I error: H_0 is incorrectly rejected; i.e., H_0 is correct but rejected

* Type II error: H_0 is incorrectly accepted i.e., H_0 is wrong but NOT rejected
- p-value: a special test statistic with a default level α rejection region $[0, \alpha]$

- Necessary components in reporting a testing result

    1. Hypotheses
    2. Name of approach
    3. Level α
    4. (Value of test statistic AND rejection region) OR p-value
    5. Conclusion: e.g., at the α level, we reject/do not reject H_0, i.e., we believe that...

Likelihood ratio test (LRT)

- Minimize the type II error rate subject to a capped type I error rate (under certain classical circumstances)

- Test statistic

$$\lambda(\boldsymbol{x}) = \frac{L(\hat{\boldsymbol{\theta}}_0; \boldsymbol{x})}{L(\hat{\boldsymbol{\theta}}; \boldsymbol{x})}$$

    - $\boldsymbol{x}$: all the observations
    - L: the likelihood function
    - $\boldsymbol{\theta}$: the unknown parameter(s)
    - $\hat{\boldsymbol{\theta}}_0$: ML estimator for $\boldsymbol{\theta}$ under H_0
    - $\hat{\boldsymbol{\theta}}$: ML estimator for $\boldsymbol{\theta}$

- (Asymptotic) rejection region

$$R_\alpha = \{\boldsymbol{x} : -2\ln\lambda(\boldsymbol{x}) \geq \chi^2_{\nu, 1-\alpha}\}$$

    - I.e., reject H_0 when $-2\ln\lambda(\boldsymbol{x}) \geq \chi^2_{\nu, 1-\alpha}$
    - $\chi^2_{\nu, 1-\alpha}$ is the $(1-\alpha)$-quantile of $\chi^2(\nu)$
    - ν: the difference in numbers of free parameters between H_0 and H_1
- (Asymptotic) p-value

$$p(\boldsymbol{x}) = 1 - F_{\chi^2(\nu)}\{-2\ln\lambda(\boldsymbol{x})\}$$

    - $F_{\chi^2(\nu)}(\cdot)$ is the cdf of $\chi^2(\nu)$